

<b>Name: Buenvenida, Ken Benedict</b>	<b>Date Performed: 10-30-2023</b>
<b>Course/Section: CPE232-CPE31S4</b>	<b>Date Submitted: 10-30-2023</b>
<b>Instructor: Engr. Jonathan Taylar</b>	<b>Semester and SY: 1st Semester 2023-2024</b>
<b>Activity 10: Install, Configure, and Manage Log Monitoring tools</b>	
<b>1. Objectives</b>	
Create and design a workflow that installs, configure and manage enterprise log monitoring tools using Ansible as an Infrastructure as Code (IaC) tool.	
<b>2. Discussion</b>	
<p>Log monitoring software scans and monitors log files generated by servers, applications, and networks. By detecting and alerting users to patterns in these log files, log monitoring software helps solve performance and security issues. System administrators use log monitoring software to detect common important events indicated by log files.</p> <p>Log monitoring software helps maintain IT infrastructure performance and pinpoints issues to prevent downtime and mitigate risks. These tools will often integrate with IT alerting software, log analysis software, and other IT issue resolution products to more aptly flesh out the IT infrastructure maintenance ecosystem.</p> <p>To qualify for inclusion in the Log Monitoring category, a product must:</p> <ul style="list-style-type: none"> <li>• Monitor the log files generated by servers, applications, or networks</li> <li>• Alert users when important events are detected</li> <li>• Provide reporting capabilities for log files</li> </ul> <p><b>Elastic Stack</b></p> <p>ELK suite stands for Elasticsearch, Kibana, Beats, and Logstash (also known as the ELK Stack). Source: <a href="https://www.elastic.co/elastic-stack">https://www.elastic.co/elastic-stack</a></p> <p>The Elastic Stack is a group of open source products from Elastic designed to help users take data from any type of source and in any format, and search, analyze and visualize that data in real time. The product group was formerly known as the ELK Stack for the core products in the group -- Elasticsearch, Logstash and Kibana -- but has been rebranded as the Elastic Stack. A fourth product, Beats, was subsequently added to the stack. The Elastic Stack can be deployed on premises or made available as software as a service (SaaS). Elasticsearch supports Amazon Web Services (AWS), Google Cloud Platform and Microsoft Azure.</p>	

## GrayLog

Graylog is a powerful platform that allows for easy log management of both structured and unstructured data along with debugging applications.

It is based on Elasticsearch, MongoDB, and Scala. Graylog has a main server, which receives data from its clients installed on different servers, and a web interface, which visualizes the data and allows to work with logs aggregated by the main server.

We use Graylog primarily as the stash for the logs of the web applications we build. However, it is also effective when working with raw strings (i.e. syslog): the tool parses it into the structured data we need. It also allows advanced custom search in the logs using structured queries. In other words, when integrated properly with a web app, Graylog helps engineers to analyze the system behavior on almost per code line basis.

Source: <https://www.graylog.org/products/open-source>

### 3. Tasks

1. Create a playbook that:
  - a. Install and configure Elastic Stack in separate hosts (Elastic Search, Kibana, Logstash)
2. Apply the concept of creating roles.
3. Describe how you did step 1. (Provide screenshots and explanations in your report. Make your report detailed such that it will look like a manual.)
4. Show an output of the installed Elastic Stack for both Ubuntu and CentOS.
5. Make sure to create a new repository in GitHub for this activity.

### 4. Output (screenshots and explanations)

#### Step 1: Pull or Clone the Repository that you just created

```
ken@controlNode:~$ git clone git@github.com:KBDBuenvenida/HOA10.git
Cloning into 'HOA10'...
remote: Enumerating objects: 3, done.
remote: Counting objects: 100% (3/3), done.
remote: Total 3 (delta 0), reused 0 (delta 0), pack-reused 0
Receiving objects: 100% (3/3), done.
```

#### Step 2: Create a Playbook named 'elastic.yml'

```
ken@controlNode:~$ cd HOA10
ken@controlNode:~/HOA10$ sudo nano elastic.yml
[sudo] password for ken:
```

**Step 3: Create an ansible.cfg or copy the old ansible.cfg in your other directories**

```
ken@controlNode:~/HOA9$ sudo cp ansible.cfg ~/HOA10
```

**Step 4: Check if the ansible.cfg is in HOA10 directory**

```
ken@controlNode:~$ cd HOA10
ken@controlNode:~/HOA10$ ls
ansible.cfg  elastic.yml  README.md
```

**Step 5: Create or copy the Inventory file**

```
ken@controlNode:~/HOA9$ sudo cp inventory ~/HOA10
```

**Step 6: Check if the inventory file is in HOA10 directory**

```
ken@controlNode:~/HOA9$ sudo cp inventory ~/HOA10
```

**Step 7: Create a directory named 'roles'**

```
ken@controlNode:~/HOA10$ mkdir roles
```

**Step 8: Create a directory inside roles and name it as elasticsearch, kibana, and logstash**

```
ken@controlNode:~/HOA10/roles$ mkdir elasticsearch kibana logstash
ken@controlNode:~/HOA10/roles$ ls
elasticsearch  kibana  logstash
```

### Step 9: Create a directory inside each roles named 'tasks'

```
ken@controlNode:~/H0A10/roles/elasticsearch$ mkdir tasks
ken@controlNode:~/H0A10/roles/elasticsearch$ ls
tasks
```

```
ken@controlNode:~/H0A10/roles/kibana$ mkdir tasks
ken@controlNode:~/H0A10/roles/kibana$ ls
tasks
```

```
ken@controlNode:~/H0A10/roles/logstash$ mkdir tasks
ken@controlNode:~/H0A10/roles/logstash$ ls
tasks
```

### Step 10: Create a main.yml for Ubuntu

INPUT	<pre>ken@controlNode:~/H0A10/roles/logstash/tasks\$ sudo nano main.yml</pre>
PROCESS	<pre>--- - name: Install Logstash in CentOS   yum:     name: logstash     state: present   when: ansible_distribution == "CentOS"  - name: Install Logstash in Ubuntu   apt:     name: logstash     state: present   when: ansible_distribution == "Ubuntu"  - name: Start Logstash   service:     name: logstash     state: started     enabled: yes</pre>

### Step 11: Create a main.yml for CentOS

INPUT	<pre>ken@controlNode:~/H0A10/roles/CentOS/tasks\$ sudo nano main.yml [sudo] password for ken:</pre>
-------	---

<b>PROCESS</b>	<pre>--- - name: Enable and start Elasticsearch service   service:     name: elasticsearch     enabled: yes     state: started  - name: Enable and start Kibana Service   service:     name: kibana     enabled: true     state: restarted  - name: Enable and Start Logstash service   service:     name: logstash     enabled: yes     state: started</pre>
<b>Step 12: Test the playbook</b>	
<b>INPUT</b>	<pre>ken@controlNode:~/HOA10\$ sudo nano Elastic.yml [sudo] password for ken: </pre>

PROCESS	<pre> - -- - hosts: all   become: true   pre_tasks:      - name: Install updates (CentOS)       yum:         update_only: yes         update_cache: yes         when: ansible_distribution == "CentOS"      - name: Install updates (Ubuntu)       apt:         upgrade: dist         update_cache: yes         when: ansible_distribution == "Ubuntu"  - hosts: Ubuntu   become: true   roles:     - Ubuntu  - hosts: CentOS   become: true   roles:     - CentOS </pre>
OUTPUT	<pre> ken@controlNode:~/H0A10\$ ansible-playbook --ask-become-pass Elastic.yml BECOME password:  PLAY [all] *****  TASK [Gathering Facts] ***** ok: [10.0.2.15] ok: [192.168.56.102]  TASK [Install updates (CentOS)] ***** skipping: [192.168.56.102] skipping: [10.0.2.15]  TASK [Install updates (Ubuntu)] ***** ok: [10.0.2.15] ok: [192.168.56.102]  PLAY [Ubuntu] *****  TASK [Gathering Facts] ***** ok: [192.168.56.102]  TASK [Ubuntu : Update apt cache] ***** changed: [192.168.56.102] </pre>

```

TASK [Ubuntu : Update apt cache] *****
changed: [192.168.56.102]

TASK [Ubuntu : Install Java] *****
ok: [192.168.56.102]

TASK [Ubuntu : Enable and Start ElasticSearch service] *****
ok: [192.168.56.102]

TASK [Ubuntu : Start Elasticsearch service] *****
ok: [192.168.56.102]

TASK [Ubuntu : Enable Kibana] *****
ok: [192.168.56.102]

TASK [Ubuntu : Start Kibana] *****
ok: [192.168.56.102]

TASK [Ubuntu : Enable and start Logstash service] *****
ok: [192.168.56.102]

PLAY [CentOS] *****

TASK [Gathering Facts] *****
ok: [192.168.56.106]

TASK [CentOS : Enable and start ElasticSearch service] *****
ok: [192.168.56.106]

TASK [CentOS : Enable and start Kibana Service] *****
changed: [192.168.56.106]

TASK [CentOS : Enable and Start Logstash service] *****
ok: [192.168.56.106]

PLAY RECAP *****
192.168.56.102      : ok=10    changed=1    unreachable=0    failed=0    s
kipped=1    rescued=0    ignored=0
192.168.56.106      : ok=6     changed=1    unreachable=0    failed=0    s
kipped=1    rescued=0    ignored=0

```

**Step 13: Test if Elasticsearch, Kibana, and Logstash is active**

**Elasticsearch**

```
ken@controlNode2:~$ systemctl status elasticsearch
● elasticsearch.service - Elasticsearch
   Loaded: loaded (/lib/systemd/system/elasticsearch.service; enabled; vendor
   Active: active (running) since Sun 2023-10-29 22:44:27 PST; 3h 56min ago
     Docs: https://www.elastic.co
   Main PID: 901 (java)
    Tasks: 65 (limit: 6896)
   Memory: 3.2G
      CPU: 1min 42.095s
   CGroup: /system.slice/elasticsearch.service
           └─ 901 /usr/share/elasticsearch/jdk/bin/java -Xms4m -Xmx64m -XX:+Use
              2853 /usr/share/elasticsearch/jdk/bin/java -Des.networkaddress.c
              3447 /usr/share/elasticsearch/modules/x-pack-ml/platform/linux-x86

Oct 29 22:36:49 controlNode2 systemd[1]: Starting Elasticsearch...
Oct 29 22:39:19 controlNode2 systemd-entrypoint[901]: Oct 29, 2023 10:39:19 PM
Oct 29 22:39:19 controlNode2 systemd-entrypoint[901]: WARNING: COMPAT locale pr
Oct 29 22:44:27 controlNode2 systemd[1]: Started Elasticsearch.
lines 1-17/17 (END)
```

```
[ken@localhost ~]$ systemctl status elasticsearch
● elasticsearch.service - Elasticsearch
   Loaded: loaded (/usr/lib/systemd/system/elasticsearch.service; enabled; vend
   Active: active (running) since Mon 2023-10-30 02:37:57 PST; 57s ago
     Docs: https://www.elastic.co
   Main PID: 2982 (java)
    Tasks: 82 (limit: 23004)
   Memory: 1.6G
   CGroup: /system.slice/elasticsearch.service
           └─ 2982 /usr/share/elasticsearch/jdk/bin/java -Xms4m -Xmx64m -XX:+Use
              3075 /usr/share/elasticsearch/jdk/bin/java -Des.networkaddress.cac
              3134 /usr/share/elasticsearch/modules/x-pack-ml/platform/linux-x86

Oct 30 02:29:49 localhost.localdomain systemd[1]: Starting Elasticsearch...
Oct 30 02:32:55 localhost.localdomain systemd-entrypoint[2982]: Oct 30, 2023 2:
Oct 30 02:32:55 localhost.localdomain systemd-entrypoint[2982]: WARNING: COMPAT
Oct 30 02:37:57 localhost.localdomain systemd[1]: Started Elasticsearch.
lines 1-16/16 (END)
```



## Kibana

```
ken@controlNode2:~$ systemctl status kibana
● kibana.service - Kibana
   Loaded: loaded (/lib/systemd/system/kibana.service; enabled; vendor preset:
   Active: active (running) since Sun 2023-10-29 22:23:55 PST; 5s ago
     Docs: https://www.elastic.co
   Main PID: 5231 (node)
    Tasks: 7 (limit: 4599)
   Memory: 43.2M
      CPU: 709ms
   CGroup: /system.slice/kibana.service
           └─5231 /usr/share/kibana/bin/../../node/bin/node /usr/share/kibana/bi>

Oct 29 22:23:55 controlNode2 systemd[1]: Started Kibana.
Oct 29 22:23:56 controlNode2 kibana[5231]: Kibana is currently running with leg>
lines 1-13/13 (END)
[2]+  Stopped                  systemctl status kibana
^Z
```

```
[ken@localhost ~]$ systemctl status kibana
● kibana.service - Kibana
   Loaded: loaded (/usr/lib/systemd/system/kibana.service; enabled; vendor pres>
   Active: active (running) since Sun 2023-10-29 22:34:03 PST; 4s ago
     Docs: https://www.elastic.co
   Main PID: 45175 (node)
    Tasks: 7 (limit: 10927)
   Memory: 16.0M
   CGroup: /system.slice/kibana.service
           └─45175 /usr/share/kibana/bin/../../node/bin/node /usr/share/kibana/bin>

Oct 29 22:34:03 localhost.localdomain systemd[1]: Started Kibana.
Oct 29 22:34:06 localhost.localdomain kibana[45175]: Kibana is currently runnin>
lines 1-12/12 (END)
```

## Logstash

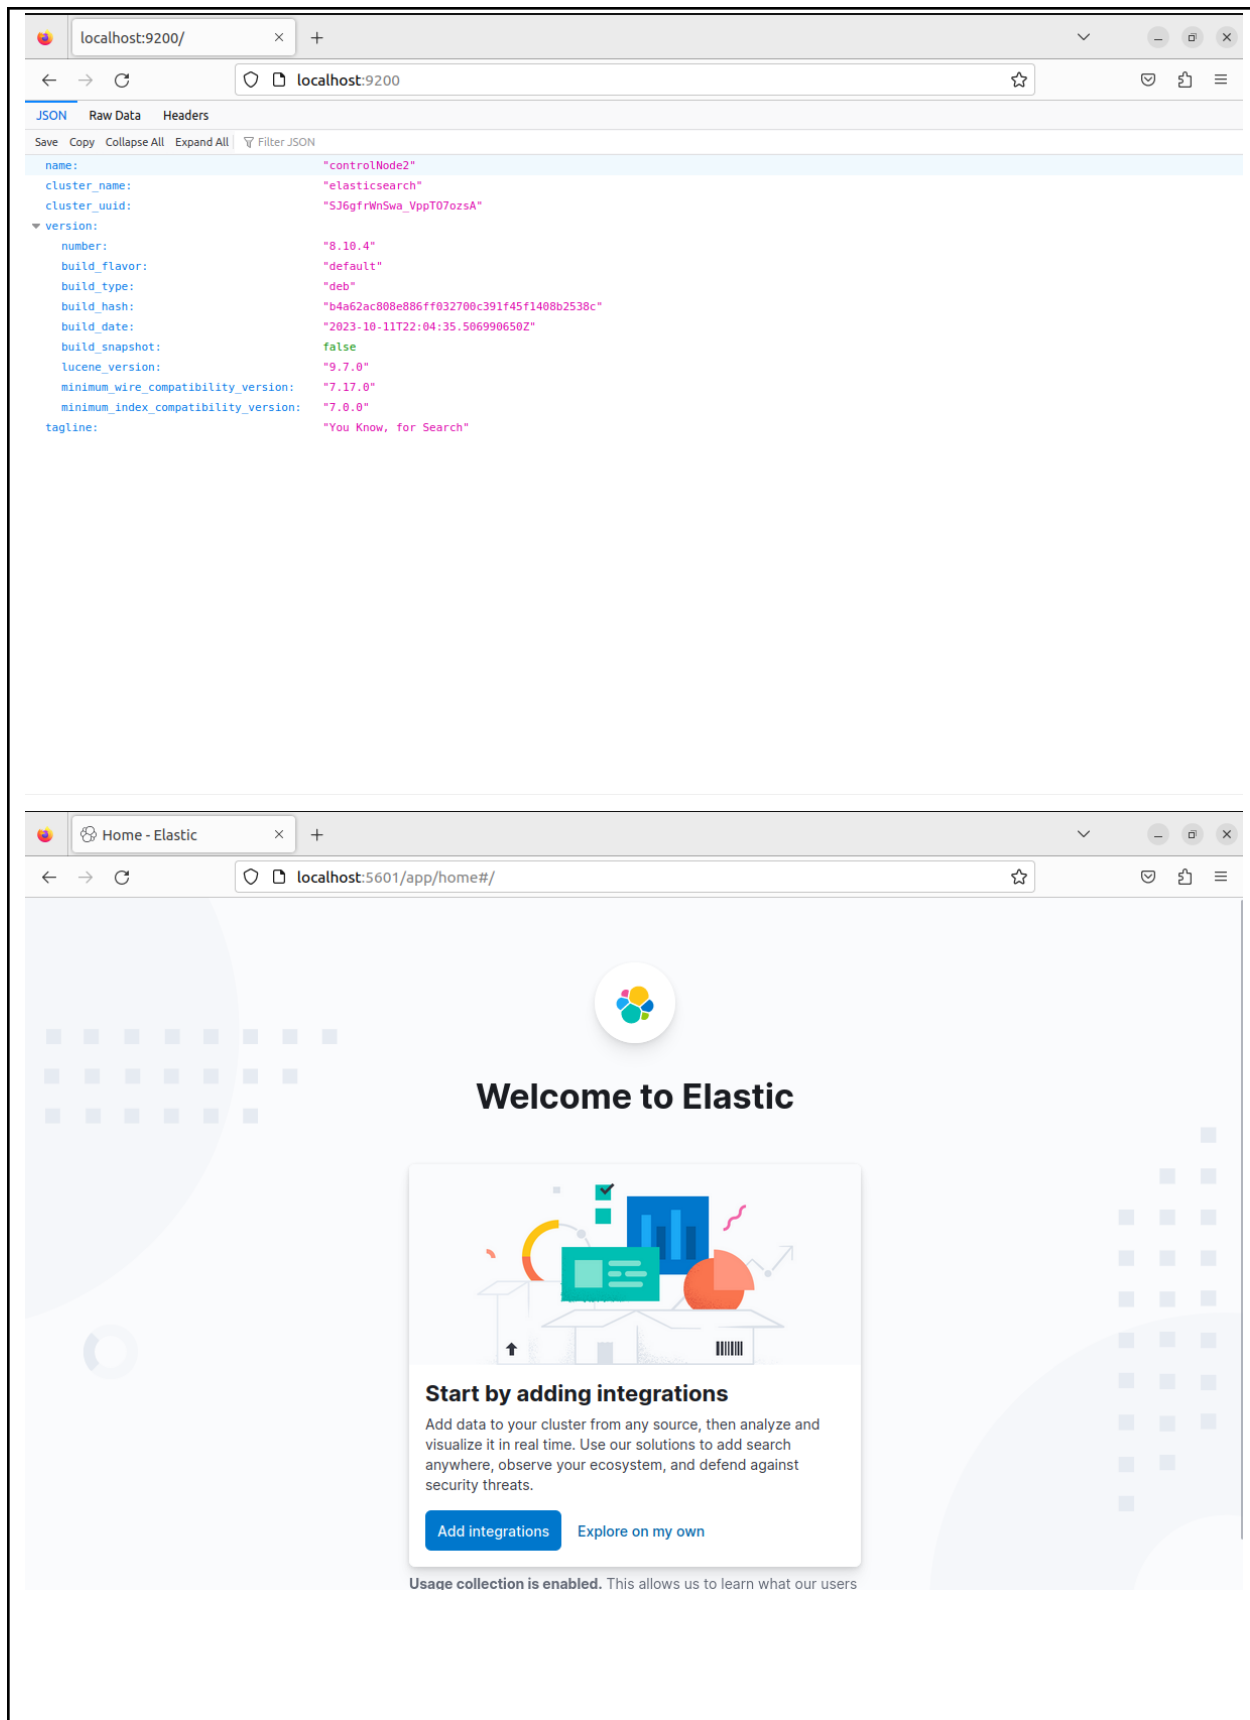
```
ken@controlNode2:~$ systemctl status logstash
● logstash.service - logstash
   Loaded: loaded (/lib/systemd/system/logstash.service; enabled; vendor preset: enabled)
   Active: active (running) since Sun 2023-10-29 22:27:51 PST; 7s ago
     Main PID: 5473 (java)
        Tasks: 14 (limit: 4599)
      Memory: 150.5M
         CPU: 2.518s
    CGroup: /system.slice/logstash.service
            └─5473 /usr/share/logstash/jdk/bin/java -Xms1g -Xmx1g -Djava.awt.headless

Oct 29 22:27:51 controlNode2 systemd[1]: Started logstash.
Oct 29 22:27:55 controlNode2 logstash[5473]: Using bundled JDK: /usr/share/logstash/jdk
lines 1-12/12 (END)
```

```
[ken@localhost ~]$ sudo systemctl enable logstash
Created symlink /etc/systemd/system/multi-user.target.wants/logstash.service → /usr/lib/systemd/system/logstash.service.
[ken@localhost ~]$ sudo systemctl start logstash
[ken@localhost ~]$ systemctl status logstash
● logstash.service - logstash
   Loaded: loaded (/usr/lib/systemd/system/logstash.service; enabled; vendor preset: enabled)
   Active: active (running) since Mon 2023-10-30 02:05:35 PST; 10s ago
     Main PID: 9091 (java)
        Tasks: 14 (limit: 10927)
      Memory: 172.8M
         CPU: 0.000s
    CGroup: /system.slice/logstash.service
            └─9091 /usr/share/logstash/jdk/bin/java -Xms1g -Xmx1g -Djava.awt.headless

Oct 30 02:05:35 localhost.localdomain systemd[1]: Started logstash.
Oct 30 02:05:35 localhost.localdomain logstash[9091]: Using bundled JDK: /usr/share/logstash/jdk
lines 1-11/11 (END)
```

**Step 14: Check if its working in the web browser**  
**Ubuntu**



# CentOS

Activities Firefox Oct 30 06:09

localhost:9200/

localhost:9200

Import bookmarks... Centos Wiki Documentation Forums

JSON Raw Data Headers

Save Copy Collapse All Expand All Filter JSON


```
name: "localhost.localdomain"
cluster_name: "elasticsearch"
cluster_uuid: "4MdhLaK8ShiCR1qOmABYLA"
version:
  number: "8.10.4"
  build_flavor: "default"
  build_type: "rpm"
  build_hash: "b4a62ac808e886ff032700c391f45f1408b2538c"
  build_date: "2023-10-11T22:04:35.506990650Z"
  build_snapshot: false
  lucene_version: "9.7.0"
  minimum_wire_compatibility_version: "7.17.0"
  minimum_index_compatibility_version: "7.0.0"
tagline: "You Know, for Search"
```

Activities Firefox Oct 30 06:10

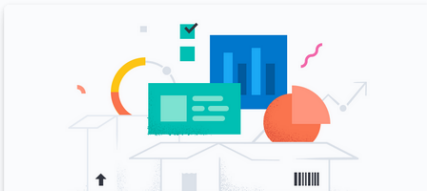
Home - Elastic

localhost:5601/app/home#/

Import bookmarks... Centos Wiki Documentation Forums



## Welcome to Elastic



### Start by adding integrations

Add data to your cluster from any source, then analyze and visualize it in real time. Use our solutions to add search anywhere, observe your ecosystem, and defend against security threats.

[Add integrations](#) [Explore on my own](#)

5. Make sure to create a new repository in GitHub for this activity.


### Step 15: Create a Repository named 'HOA10'

## Create a new repository

A repository contains all project files, including the revision history. Already have a project repository elsewhere? [Import a repository.](#)

*Required fields are marked with an asterisk (\*).*

Owner \*

 KBDBuenvenida ▾

/


Repository name \*

HOA10


✔ HOA10 is available.

Great repository names are short and memorable. Need inspiration? How about [symmetrical-telegram](#) ?

Description (optional)

☒  **Public**

Anyone on the internet can see this repository. You choose who can commit.

☐  **Private**

You choose who can see and commit to this repository.

**Initialize this repository with:**

☒ **Add a README file**

This is where you can write a long description for your project. [Learn more about READMEs.](#)

**Add .gitignore**

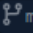
.gitignore template: **None** ▾


Choose which files not to track from a list of templates. [Learn more about ignoring files.](#)

**Choose a license**

License: **None** ▾

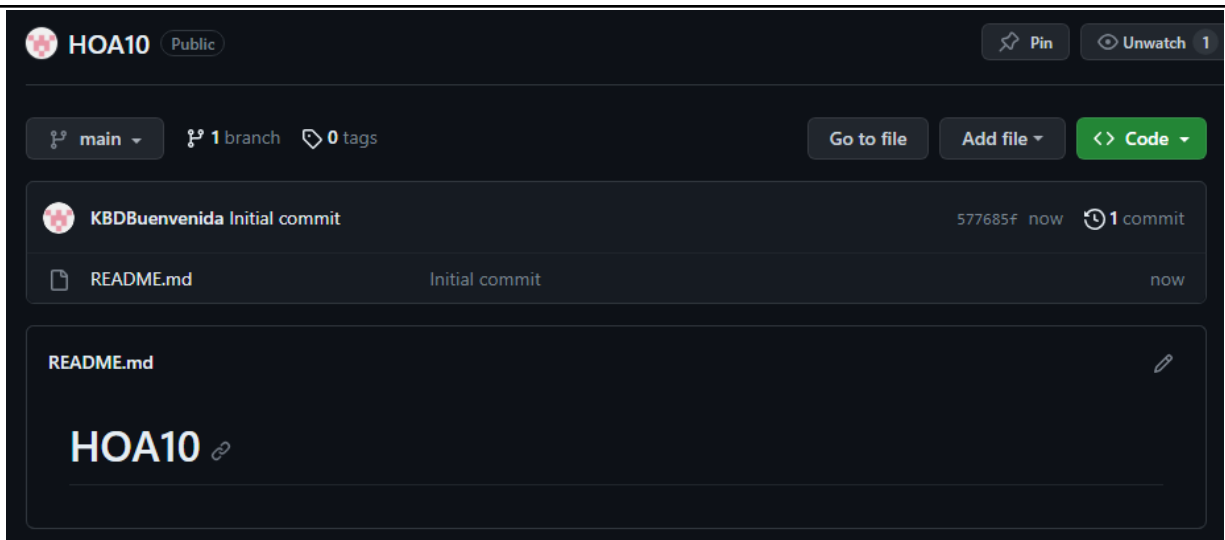
A license tells others what they can and can't do with your code. [Learn more about licenses.](#)

This will set  **main** as the default branch. Change the default name in your [settings](#).

 You are creating a public repository in your personal account.

Creating repository...

### Step 16: Check if the new repository is created



#### Step 17: Git add \*

```
ken@controlNode:~/HOA10$ git add *
```

#### Step 18: Git commit -m "HOA10"

```
ken@controlNode:~/HOA10$ git commit -m "HOA10"
[main b64a5b0] HOA10
6 files changed, 144 insertions(+)
create mode 100644 Elastic.yml
create mode 100644 ansible.cfg
create mode 100644 inventory
create mode 100644 roles/CentOS/tasks/main.yml
create mode 100644 roles/Ubuntu/main.yml
create mode 100644 roles/Ubuntu/tasks/main.yml
```

#### Step 19: Git push origin

```
ken@controlNode:~/HOA10$ git push origin
Enumerating objects: 14, done.
Counting objects: 100% (14/14), done.
Compressing objects: 100% (10/10), done.
Writing objects: 100% (13/13), 1.62 KiB | 1.62 MiB/s, done.
Total 13 (delta 1), reused 0 (delta 0), pack-reused 0
remote: Resolving deltas: 100% (1/1), done.
To github.com:KBDBuenvenida/HOA10.git
577685f..b64a5b0 main -> main
```

#### Reflections:

Answer the following:

1. What are the benefits of having log monitoring tool?

- **ELK Stack can be used to monitor performance and health of the user's system and application. The user can use it to track CPU usage, memory usage, and response times. It can also be used to provide alerts when errors or problems are detected. ELK Stack can also be used to analyze the user's business data to gain insights to customers, products, and operations. ELK Stack is a good choice for organizations that are looking to improve their log management and analytics capabilities.**

**Conclusions:**

- **What I did in this activity is to manually install ELK Stack since everytime I try to install it using a playbook it displays an error or it corrupts my virtual machine. In conclusion, I learned a lot about ELK Stack in this activity and that it can be used to monitor my machine performance and health. It was stressful to do at first but once I was able to start the ELK Stack it was smooth after that, I was able to use and setup elasticsearch, kibana, and logstash. I was able to use ELK Stack in a browser where I can use a lot of their features.**